


---

# SENTIMENT ANALYSIS OF IMDB REVIEWS

---



Advanced Text Analytics Project  
Merve Pakcan



# TABLE OF CONTENTS

**01** ABSTRACT & OBJECTIVES

**02** DATASET DESCRIPTION

**03** DATA PREPROCESSING

**04** MODEL STRUCTURES &  
IMPLEMENTATION

Logistic Regression, Naive Bayes, Neural  
Network

**05** HYPERPARAMETER TUNING  
& MODIFICATIONS

**06** SENTIMENT ANALYSIS IN  
SAS

**07** RESULTS & CONCLUSION



# INTRODUCTION

User-generated content is growing. Movie reviews are a great source to understand audience sentiment.

## Objective:

To compare the performance of three machine learning models **Naive Bayes**, **Logistic Regression**, and a **Neural Network** for classifying IMDB movie reviews as **positive** or **negative**.

## What is Sentiment Analysis?

A subfield of NLP, detects and classifies **emotions** or **opinions** in text.



# DATASET DESCRIPTION

"IMDB Dataset of 50K Movie Reviews", available on **Kaggle**

**50,000 REVIEWS, 2 COLUMNS**

80% training, 20% testing

**REVIEW**

**SENTIMENT**

**Source of the data:**

**Name:** Large Movie Review Dataset (IMDB)

**Source:** Stanford AI Lab

**Authors:** Andrew L. Maas et al., ACL 2011

**Size:** 50,000 reviews (25K train / 25K test)

**Balanced** dataset (50% positive, 50% negative)



## DATA PREPROCESSING

### Naive Bayes & Logistic Regression

- Lowercasing
- Remove HTML, URLs, punctuation
- Tokenization (TweetTokenizer)
- Stopword removal
- Stemming (Porter)
- (NB only) Frequency dictionary creation

### Neural Network

- Tokenization (Keras Tokenizer, vocab=10,000)
- Convert to sequences
- Padding to fixed length
- No stemming or stopword removal
- Word representation via Embedding Layer

### Why Different Preprocessing?



# MODEL STRUCTURES & IMPLEMENTATION

## 1. LOGISTIC REGRESSION

A supervised binary classifier to predict sentiment(positive/negative),


**Built with Scikit-learn's LogisticRegression (default parameters)**

**Input:** Preprocessed text converted into numeric features

**Output:** Probability mapped to sentiment label (0 or 1)

**Initial test accuracy: 62.08%**

- Model predicted **positive** for mixed/neutral reviews
- "I expected a great movie, but it was boring and full of clichés." → **Positive (0.62)**
- **Biased toward positive words**



- **Recall imbalance:** Positive = 0.94, Negative = 0.30  
(3488 was negative but model says positive)

## 2. NAIVE BAYES

A probabilistic classifier that applies Bayes' Theorem, assuming word independence.

Built with NLTK's NaiveBayesClassifier

**Input:** Cleaned, stemmed tokens with (word, label) frequency pairs

**Output:** Class label based on word-level probability estimates

**Initial test accuracy: 85.0%**

**Balanced recall** for both classes ( $\approx 0.85$ )


Correctly predicted:

- 4,249 negative, 4,217 positive

**Strength:** Effective on clear sentiment

*The movie had some good moments, but overall it felt flat.*

**Limitation:** Tended to label mixed-tone reviews as negative





### 3. NEURAL NETWORK

A neural model built using embedding and dense layers to classify review sentiment.

**Built with:** Keras Sequential model

**Input:** Padded token sequences

**Output:** Sentiment label (0 or 1) via sigmoid

**Test Accuracy: 88.00%**

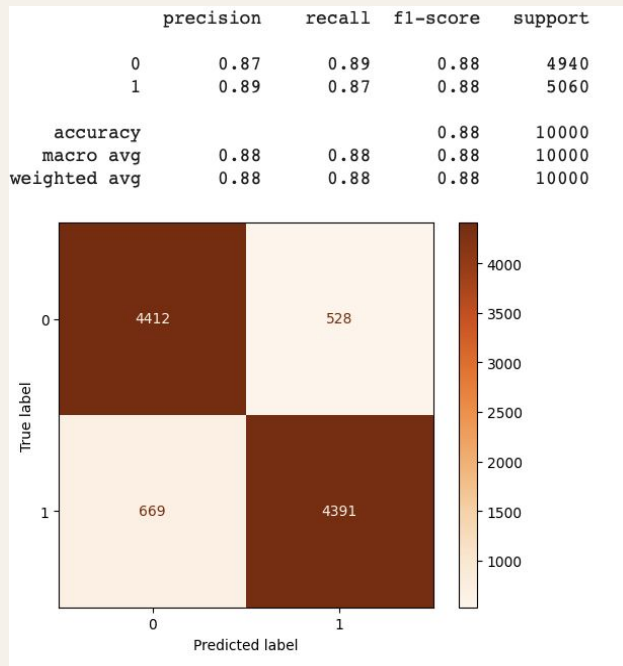
Balanced **precision & recall**: both **0.88**

Correctly classified **neutral/mixed** reviews

“The movie had good moments, but felt flat.” → **Negative (0.33)**

Model learned word patterns through **embedding layer**

Used **EarlyStopping** to avoid overfitting and boost generalization







# SENTIMENT ANALYSIS IN SAS

Pipeline completed successfully.

- **Text Parsing:** 64,477 unique terms extracted with word types ( noun, verb)
- **Sentiment Analysis:** Automatically labeled 49,999 reviews as *positive*, *negative*, or *neutral*
- **Topic Modeling:** Clustered reviews into key review topics
- **Categories:** Organized output into clear, interpretable groups



## Hyperparameter Tuning Logistic Regression

Best Params: {'C': 1, 'penalty': 'l2', 'solver': 'liblinear'}  
Best CV Accuracy: 0.8862971683960496

	precision	recall	f1-score	support
0	0.90	0.87	0.88	4983
1	0.88	0.90	0.89	5017
accuracy			0.89	10000
macro avg	0.89	0.88	0.88	10000
weighted avg	0.89	0.89	0.88	10000

**Method:** GridSearchCV (CV=5, scoring='accuracy')

**Parameters Tuned:** C, penalty, solver

**Before tuning:** 62% accuracy, low recall on negative class

**After tuning: Accuracy: 88.0%, Recall (Neg):** ↑ from 0.30 to 0.87,

**Recall (Pos):** stayed high at 0.90

Model became more balanced and consistent

## Naive Bayes

Already in optimum

Before: 85% After: 85%

Best alpha: {'alpha': 10}

Best CV accuracy: 0.8546963682960371

	precision	recall	f1-score	support
0	0.84	0.85	0.85	4983
1	0.85	0.84	0.85	5017
accuracy			0.85	10000
macro avg	0.85	0.85	0.85	10000
weighted avg	0.85	0.85	0.85	10000

## Neural Network

88% accuracy

Tuned **vocab size**((10k), **max length**(200), and **embedding dim**(16)

Added **EarlyStopping** to prevent overfitting

Trained with **batch data + validation set**

Tested on **ambiguous/neutral reviews**

Bidirectional LSTM and lower LR (Adam 0.0005) —  
no accuracy gain (still 88%).



# CONCLUSION

- Compared **Logistic Regression**, **Naive Bayes**, and **Neural Network** on IMDB reviews
- **Naive Bayes**: 85% on clear sentiment, weak on mixed
- **Logistic Regression**: improved from 62% → 88% after tuning
- **Neural Network**: highest accuracy (88%), best on ambiguous reviews
- **SAS**: verified sentiment and topic labeling on full dataset



# THANKS

Do you have any questions?

